

# **SURV727\_assignment5**

Zihan Tong

GitHub: [https://github.com/tzihan/SURV727\\_assignment5](https://github.com/tzihan/SURV727_assignment5)

You may work in pairs or individually for this assignment. Make sure you join a group in Canvas if you are working in pairs. Turn in this assignment as an HTML or PDF file to ELMS. Make sure to include the R Markdown or Quarto file that was used to generate it. Include the GitHub link for the repository containing these files.

```
library(censusapi)
library(tidyverse)

-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr     1.1.4      v readr     2.1.5
v forcats   1.0.0      v stringr   1.5.2
v ggplot2   3.5.2      v tibble    3.3.0
v lubridate 1.9.4      v tidyr    1.3.1
v purrr     1.1.0

-- Conflicts -----
x dplyr::filter() masks stats::filter()
x dplyr::lag()    masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become non-conflicting.

library(magrittr)
```

Attaching package: 'magrittr'

The following object is masked from 'package:purrr':

```
set_names
```

The following object is masked from 'package:tidyverse':

```
extract
```

```
library(factoextra)
```

Welcome! Want to learn more? See two factoextra-related books at <https://goo.gl/ve3WBa>

## Exploring ACS Data

In this notebook, we use the Census API to gather data from the American Community Survey (ACS). This requires an access key, which can be obtained here:

[https://api.census.gov/data/key\\_signup.html](https://api.census.gov/data/key_signup.html)

```
cs_key <- readLines("Census_key.txt")
```

```
acs_il_c <- getCensus(name = "acs/acss5",
                        vintage = 2016,
                        vars = c("NAME",
                                "B01003_001E",
                                "B19013_001E",
                                "B19301_001E"),
                        region = "county:*",
                        regionin = "state:17",
                        key = cs_key) %>%
  rename(pop = B01003_001E,
         hh_income = B19013_001E,
         income = B19301_001E)
head(acs_il_c)
```

|   | state | county |                            | NAME   | pop   | hh_income | income |
|---|-------|--------|----------------------------|--------|-------|-----------|--------|
| 1 | 17    | 067    | Hancock County, Illinois   | 18633  | 50077 | 25647     |        |
| 2 | 17    | 063    | Grundy County, Illinois    | 50338  | 67162 | 30232     |        |
| 3 | 17    | 091    | Kankakee County, Illinois  | 111493 | 54697 | 25111     |        |
| 4 | 17    | 043    | DuPage County, Illinois    | 930514 | 81521 | 40547     |        |
| 5 | 17    | 003    | Alexander County, Illinois | 7051   | 29071 | 16067     |        |
| 6 | 17    | 129    | Menard County, Illinois    | 12576  | 60420 | 31323     |        |

Pull map data for Illinois into a data frame.

```
il_map <- map_data("county", region = "illinois")
head(il_map)
```

|   | long      | lat      | group | order | region   | subregion |
|---|-----------|----------|-------|-------|----------|-----------|
| 1 | -91.49563 | 40.21018 | 1     | 1     | illinois | adams     |
| 2 | -90.91121 | 40.19299 | 1     | 2     | illinois | adams     |
| 3 | -90.91121 | 40.19299 | 1     | 3     | illinois | adams     |
| 4 | -90.91121 | 40.10704 | 1     | 4     | illinois | adams     |
| 5 | -90.91121 | 39.83775 | 1     | 5     | illinois | adams     |
| 6 | -90.91694 | 39.75754 | 1     | 6     | illinois | adams     |

Join the ACS data with the map data. Note that `il_map` has a column `subregion` which includes county names. We need a corresponding variable in the ACS data to join both data sets. This needs some transformations, among which the function `tolower()` might be useful. Call the joined data `acs_map`.

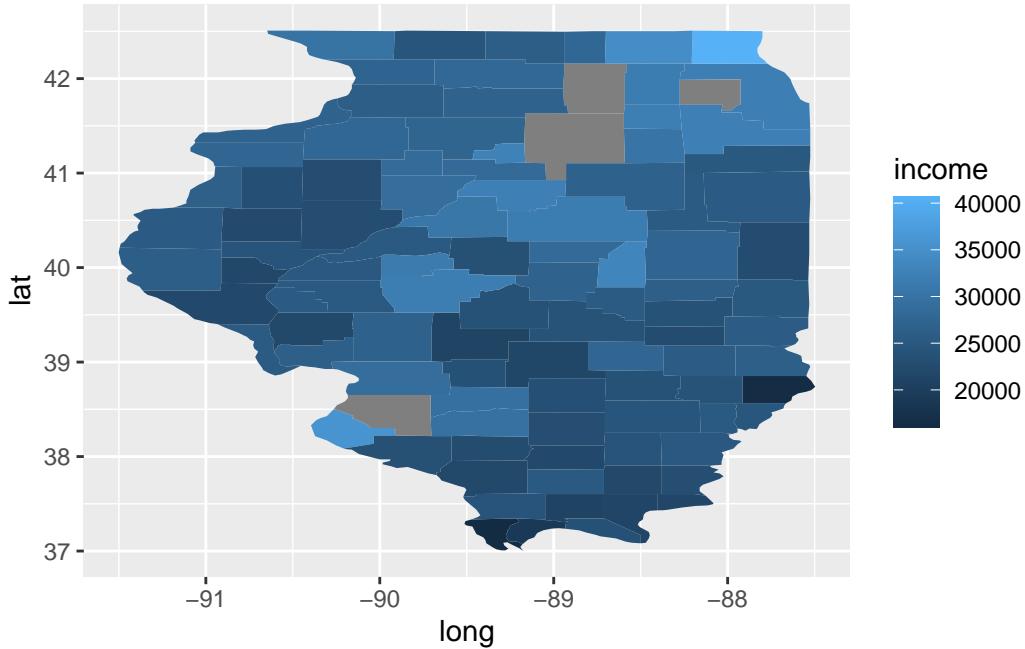
```
acs_il_c <- acs_il_c %>%
  mutate(
    county = NAME %>%
      gsub(" County, Illinois", "", .) %>%
      tolower()
  )

il_map <- il_map %>%
  mutate(subregion = tolower(subregion))

acs_map <- il_map %>%
  left_join(acs_il_c, by = c("subregion" = "county"))
```

After you do this, plot a map of Illinois with Counties colored by per capita income.

```
ggplot(acs_map) +
  geom_polygon(aes(x = long,
                  y = lat,
                  group = group,
                  fill = income))
```



## Hierarchical Clustering

We want to find clusters of counties that are similar in their population, average household income and per capita income. First, clean the data so that you have the appropriate variables to use for clustering. Next, create the distance matrix of the cleaned data. This distance matrix can be used to cluster counties, e.g. using the ward method.

```
clust_data <- acs_il_c %>%
  select(pop, hh_income, income) %>%
  drop_na()

clust_scaled <- scale(clust_data)

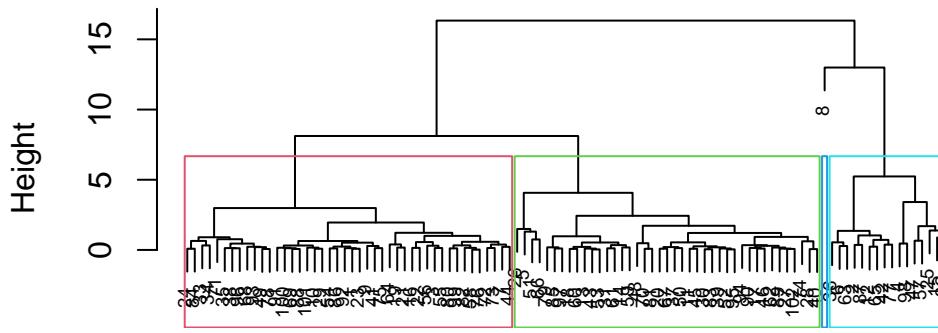
dist_matrix <- dist(clust_scaled)

hc <- hclust(dist_matrix, method = "ward.D2")
```

Plot the dendrogram to find a reasonable number of clusters. Draw boxes around the clusters of your cluster solution.

```
plot(hc, main = "Hierarchical Clustering Dendrogram",
      xlab = "Counties", ylab = "Height", cex = 0.6)
rect.hclust(hc, k = 4, border = 2:5)
```

## Hierarchical Clustering Dendrogram



Counties  
hclust (\*, "ward.D2")

Visualize the county clusters on a map. For this task, create a new `acs_map` object that now also includes cluster membership as a new column. This column should be called `cluster`.

```
clusters <- cutree(hc, k = 4)

acs_il_c$cluster <- clusters

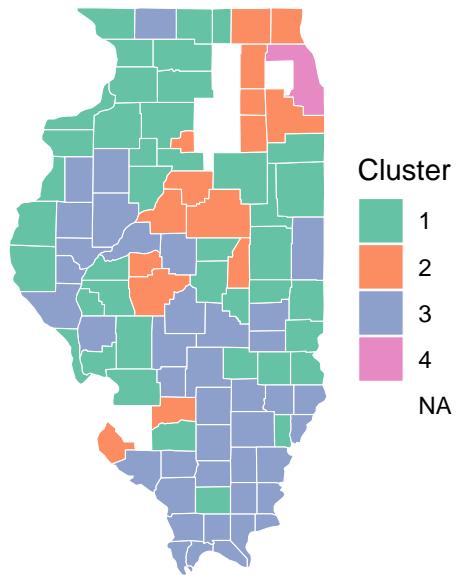
acs_map <- il_map %>%
  left_join(acs_il_c, by = c("subregion" = "county"))

ggplot(acs_map) +
  geom_polygon(aes(x = long,
                  y = lat,
                  group = group,
                  fill = factor(cluster)),
               color = "white", size = 0.2) +
  coord_fixed(1.3) +
  scale_fill_brewer(palette = "Set2", name = "Cluster") +
```

```
labs(title = "County Clusters in Illinois (Hierarchical Clustering)") +  
  theme_void()
```

Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.  
i Please use `linewidth` instead.

County Clusters in Illinois (Hierarchical Clustering)



## Census Tracts

For the next section we need ACS data on a census tract level. We use the same variables as before.

```
acs_il_t <- getCensus(name = "acs/acss5",  
                        vintage = 2016,  
                        vars = c("NAME",  
                                "B01003_001E",  
                                "B19013_001E",  
                                "B19301_001E"),  
                        region = "tract:*",  
                        regionin = "state:17",  
                        key = cs_key) %>%  
  mutate_all(fun_ifelse(.==666666666, NA, .))) %>%
```

```

    rename(pop = B01003_001E,
           hh_income = B19013_001E,
           income = B19301_001E)

Warning: `fun` was deprecated in dplyr 0.8.0.
i Please use a list of either functions or lambdas:

# Simple named list: list(mean = mean, median = median)

# Auto named with `tibble::lst()`: tibble::lst(mean, median)

# Using lambdas list(~ mean(., trim = .2), ~ median(., na.rm = TRUE))

head(acs_il_t)

```

|   | state | county | tract     | NAME  | pop  |
|---|-------|--------|-----------|---|------|
| 1 | 17    | 031    | 806002    | Census Tract 8060.02, Cook County, Illinois | 7304 |
| 2 | 17    | 031    | 806003    | Census Tract 8060.03, Cook County, Illinois | 7577 |
| 3 | 17    | 031    | 806400    | Census Tract 8064, Cook County, Illinois    | 2684 |
| 4 | 17    | 031    | 806501    | Census Tract 8065.01, Cook County, Illinois | 2590 |
| 5 | 17    | 031    | 750600    | Census Tract 7506, Cook County, Illinois    | 3594 |
| 6 | 17    | 031    | 310200    | Census Tract 3102, Cook County, Illinois    | 1521 |
|   |       |        | hh_income | income                                      |      |
| 1 |       |        | 56975     | 23750                                       |      |
| 2 |       |        | 53769     | 25016                                       |      |
| 3 |       |        | 62750     | 30154                                       |      |
| 4 |       |        | 53583     | 20282                                       |      |
| 5 |       |        | 40125     | 18347                                       |      |
| 6 |       |        | 63250     | 31403                                       |      |

## k-Means

As before, clean our data for clustering census tracts based on population, average household income and per capita income.

```

tract_clust_data <- acs_il_t %>%
  select(pop, hh_income, income) %>%
  drop_na()

tract_scaled <- scale(tract_clust_data)

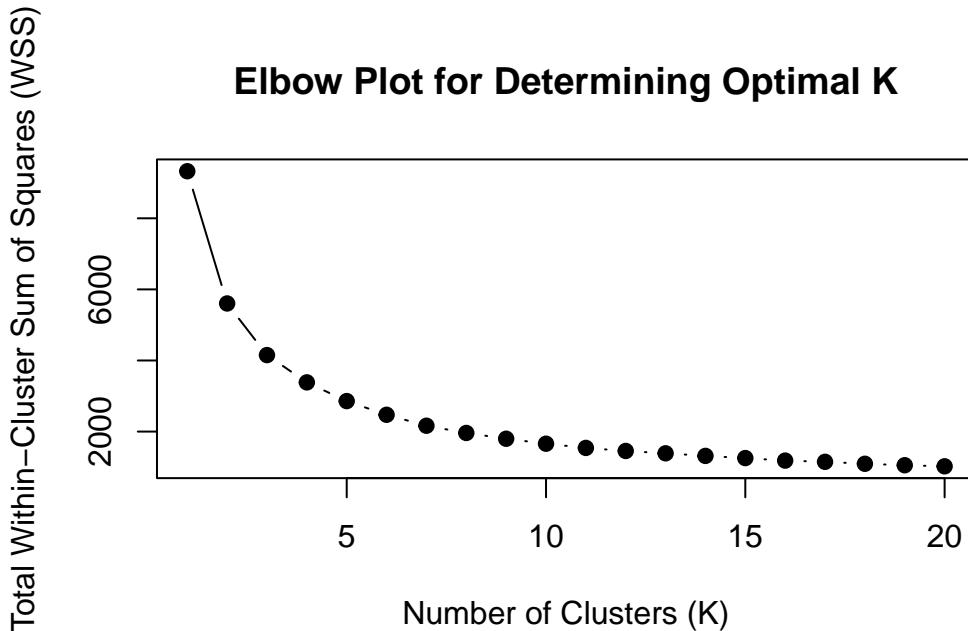
```

Since we want to use K Means in this section, we start by determining the optimal number of K that results in Clusters with low within but high between variation. Plot within cluster sums of squares for a range of K (e.g. up to 20).

```
wss_func <- function(k) {
  kmeans(tract_scaled,
    centers = k,
    nstart = 20,
    iter.max = 200)$tot.withinss
}

k_values <- 1:20
wss_values <- sapply(k_values, wss_func)

# Elbow plot
plot(k_values, wss_values,
  type = "b", pch = 19,
  xlab = "Number of Clusters (K)",
  ylab = "Total Within-Cluster Sum of Squares (WSS)",
  main = "Elbow Plot for Determining Optimal K")
```



- From the elbow plot, WSS decreases quickly up to  $K = 3$  or  $4$ , and the curve flattens afterward. Both values are reasonable, but I chose  $K = 4$  because it keeps WSS low

while providing a more detailed clustering structure.

Run `kmeans()` for the optimal number of clusters based on the plot above.

```
set.seed(123)
kmeans_result <- kmeans(
  tract_scaled,
  centers = 4,
  nstart = 20,
  iter.max = 200
)

acs_il_t$non_missing <- complete.cases(
  acs_il_t[, c("pop", "hh_income", "income")])

acs_il_t$cluster_k4 <- NA
acs_il_t$cluster_k4[acs_il_t$non_missing] <- kmeans_result$cluster
```

Find the mean population, household income and per capita income grouped by clusters. In addition, display the most frequent county that can be observed within each cluster.

```
acs_il_t$county <- sub(".*", "(.) County,.*", "\\\1", acs_il_t$NAME)

cluster_summary <- acs_il_t %>%
  filter(!is.na(cluster_k4)) %>%
  group_by(cluster_k4) %>%
  summarize(
    mean_pop = mean(pop, na.rm = TRUE),
    mean_hh_income = mean(hh_income, na.rm = TRUE),
    mean_percap_income = mean(income, na.rm = TRUE),
    .groups = "drop"
  )
cluster_summary
```

```
# A tibble: 4 x 4
  cluster_k4 mean_pop mean_hh_income mean_percap_income
    <int>     <dbl>        <dbl>            <dbl>
1         1     3764.      70344.          34486.
2         2     3137.      37770.          19719.
3         3     4134.     120509.          65518.
4         4     7011.      65178.          29126.
```

- Cluster 3 has the highest household income and per-capita income and represents the most affluent areas.
- Cluster 2 has the lowest values across all indicators and represents economically disadvantaged areas.
- Clusters 1 and 4 fall in the middle in terms of income.
- Cluster 4 has a much larger average population, likely corresponding to more urban or densely populated regions.

```
county_mode <- acs_il_t %>%
  filter(!is.na(cluster_k4)) %>%
  group_by(cluster_k4) %>%
  summarize(
    most_frequent_county = names(which.max(table(county)))
  )
county_mode
```

```
# A tibble: 4 x 2
  cluster_k4 most_frequent_county
  <int> <chr>
1       1 Cook
2       2 Cook
3       3 Cook
4       4 Cook
```

- The most frequent county in all four clusters is Cook County, likely because Cook has the largest population and the largest number of census tracts, so it dominates the clusters.

As you might have seen earlier, it's not always clear which number of clusters is the optimal choice. To automate K Means clustering, program a function based on `kmeans()` that takes K as an argument. You can fix the other arguments, e.g. such that a specific dataset is always used when calling the function.

```
run_kmeans <- function(K) {
  result <- kmeans(
    x = tract_scaled,
    centers = K,
    nstart = 20,
    iter.max = 200
  )
  return(result$cluster)
}
```

We want to utilize this function to iterate over multiple Ks (e.g.,  $K = 2, \dots, 10$ ) and – each time – add the resulting cluster membership as a new variable to our (cleaned) original data frame (`acs_il_t`). There are multiple solutions for this task, e.g. think about the `apply` family or `for` loops.

```
for (k in 2:10) {
  clusters <- run_kmeans(k)
  col_name <- paste0("cluster_k", k)
  acs_il_t[[col_name]] <- NA
  acs_il_t[[col_name]][acs_il_t$non_missing] <- clusters
}
```

Finally, display the first rows of the updated data set (with multiple cluster columns).

```
head(acs_il_t)
```

|   | state      | county     | tract       | NAME  | pop         |            |            |
|---|------------|------------|-------------|---|-------------|------------|------------|
| 1 | 17         | Cook       | 806002      | Census Tract 8060.02, Cook County, Illinois | 7304        |            |            |
| 2 | 17         | Cook       | 806003      | Census Tract 8060.03, Cook County, Illinois | 7577        |            |            |
| 3 | 17         | Cook       | 806400      | Census Tract 8064, Cook County, Illinois    | 2684        |            |            |
| 4 | 17         | Cook       | 806501      | Census Tract 8065.01, Cook County, Illinois | 2590        |            |            |
| 5 | 17         | Cook       | 750600      | Census Tract 7506, Cook County, Illinois    | 3594        |            |            |
| 6 | 17         | Cook       | 310200      | Census Tract 3102, Cook County, Illinois    | 1521        |            |            |
|   | hh_income  | income     | non_missing | cluster_k4                                  | cluster_k2  | cluster_k3 | cluster_k5 |
| 1 | 56975      | 23750      | TRUE        | 1   | 2           | 2          | 1          |
| 2 | 53769      | 25016      | TRUE        | 1   | 2           | 2          | 1          |
| 3 | 62750      | 30154      | TRUE        | 2   | 2           | 3          | 2          |
| 4 | 53583      | 20282      | TRUE        | 4   | 2           | 3          | 3          |
| 5 | 40125      | 18347      | TRUE        | 4   | 2           | 3          | 3          |
| 6 | 63250      | 31403      | TRUE        | 2   | 2           | 3          | 2          |
|   | cluster_k6 | cluster_k7 | cluster_k8  | cluster_k9                                  | cluster_k10 |            |            |
| 1 | 4          | 5          | 2           | 9   | 8           |            |            |
| 2 | 4          | 5          | 2           | 9   | 8           |            |            |
| 3 | 5          | 7          | 6           | 2   | 9           |            |            |
| 4 | 5          | 2          | 6           | 2   | 9           |            |            |
| 5 | 2          | 2          | 1           | 4   | 6           |            |            |
| 6 | 5          | 7          | 6           | 2   | 9           |            |            |